

Getting / cleaning data 2

Regular expressions

Regular expressions

We've already done some things to manipulate strings. For example, if we wanted to separate "Name" into last name and first name (including title), we could actually do that with the `separate` function:

```
titanic_train %>%  
  select(Name) %>%  
  slice(1:3) %>%  
  separate(Name, c("last_name", "first_name"), sep = ", ")  
  
##   last_name                               first_name  
## 1   Braund                               Mr. Owen Harris  
## 2 Cumings Mrs. John Bradley (Florence Briggs Thayer)  
## 3 Heikkinen                             Miss. Laina
```

Regular expressions

Notice that `separate` is looking for a regular pattern (",") and then doing something based on the location of that pattern in each string (splitting the string).

There are a variety of functions in R that can perform manipulations based on finding regular patterns in character strings.

Regular expressions

Braund, Mr. Owen Harris
Cummings, Mrs. John Bradley (Florence Briggs Thayer)
Heikkinen, Miss. Laina



, M—. pattern

Braund, Mr. Owen Harris
Cummings, Mrs. John Bradley (Florence Briggs Thayer)
Heikkinen, Miss. Laina

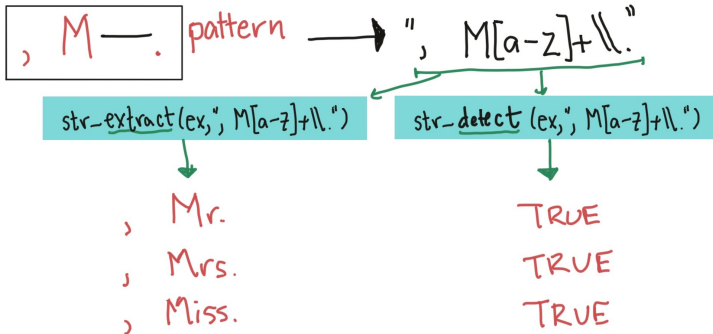
Regular expressions

Vector 'ex':

Braund, Mr. Owen Harris

Cumings, Mrs. John Bradley (Florence Briggs Thayer)

Heikkinen, Miss. Laina



Regular expressions

pattern: "Mr"

Strings

str_extract
result

str_detect
result

Mr.

Mr

TRUE

Mrs.

Mr

TRUE

Miss.

NA

FALSE

Dr.

NA

FALSE

Regular expression patterns

The easiest regular expression patterns are literal text. For example, the regular expression pattern if you're trying to match "Mr" is just "Mr":

```
ex_names <- c("Braund, Mr. Owen Harris",  
              "Cumings, Mrs. John Bradley",  
              "Heikkinen, Miss. Laina")  
str_extract(ex_names, pattern = "Mr")  
  
## [1] "Mr" "Mr" NA
```


Regular expression patterns

Regular expression patterns are case sensitive, so you won't match "Mr" with the pattern "mr":

```
ex_names <- c("Braund, Mr. Owen Harris",  
              "Cumings, Mrs. John Bradley",  
              "Heikkinen, Miss. Laina")  
str_extract(ex_names, pattern = "mr")  
  
## [1] NA NA NA
```

Regular expression patterns

There are a few characters called **metacharacters** that mean something special in regular expression patterns.

To use any of these literally in a regular expression, you need to “protect” them with two backslashes.

Regular expressions

pattern: "Mr."

<u>strings</u>	<u>str_extract result</u>	<u>str_detect result</u>
Mr.	Mr.	TRUE
Mrs.	Mrs	TRUE
Miss.	NA	FALSE
Dr.	NA	FALSE

Regular expressions

pattern: "Mr\\."

<u>Strings</u>	<u>str_extract result</u>	<u>str_detect result</u>
Mr.	Mr.	TRUE
Mrs.	NA	FALSE
Miss.	NA	FALSE
Dr.	NA	FALSE

Regular expression patterns

For example, “.” is a metacharacter, so to match “Mr.”, you need to use the pattern “Mr\\..”:

```
ex_names <- c("Braund, Mr. Owen Harris",  
              "Cumings, Mrs. John Bradley",  
              "Heikkinen, Miss. Laina")  
str_extract(ex_names, pattern = "Mr\\..")  
  
## [1] "Mr." NA    NA
```

Regular expression metacharacters

<u>Metacharacter</u>	<u>Use</u>	<u>To match literally</u>
.	match any character	"\."
*	match ≥ 0 of something	"*"
+	match ≥ 1 of something	"\+"
[]	match a character in a subset	"\[\" "]"
^	depends on context	"\^"
()	extract part of a pattern	"(\" ")
?	match zero or one of something	"\?"
{ }	customize number of times to match	"\{" "\}"
\	escape a metacharacter	"\\\""
\$	match a pattern at the end of the string	"\\$"

Regular expression patterns

pattern: "Mr[s]*\."

0 or more "s"s

Strings

str_extract
result

str_detect
result

Mr.

Mr.

TRUE

Mrs.

Mrs.

TRUE

Miss.

NA

FALSE

Dr.

NA

FALSE

Regular expression patterns

pattern: "M[a-z]+\."		
	1 or more lower case letters	
<u>strings</u>	<u>str_extract result</u>	<u>str_detect result</u>
Mr.	Mr.	TRUE
Mrs.	Mrs.	TRUE
Miss.	Miss.	TRUE
Dr.	NA	FALSE

Regular expressions

The last pattern used `[a-z]+` to match one or more lowercase letters. The `[a-z]` is a **character class**.

You can also match digits (`[0-9]`), uppercase letters (`[A-Z]`), just some letters (`[aeiou]`), etc.

You can negate a character class by starting it with `^`. For example, `[^0-9]` will match anything that **isn't** a digit.

Regular expression patterns

pattern: "[A-Z][a-z]+\."

1 uppercase character

1 or more lower case letters

Strings

str_extract
result

str_detect
result

Mr.

Mr.

TRUE

Mrs.

Mrs.

TRUE

Miss.

Miss.

TRUE

Dr.

Dr.

TRUE

Regular expressions

The `str_detect` function will look through each element of a character vector for a designated pattern. If the pattern is there, it will return `TRUE`, and otherwise `FALSE`. The convention is:

```
## Generic code
str_detect(string = [vector you want to check],
           pattern = [pattern you want to check for])
```

For example, to create a logical vector specifying which of the Titanic passenger names include "Mrs.", you can call:

```
mrs <- str_detect(titanic_train$Name, "Mrs\\.")
head(mrs)
```

```
## [1] FALSE TRUE FALSE TRUE FALSE FALSE
```

Regular expressions

The result is a logical vector, so `str_detect` can be used in `filter` to subset data to only rows where the passenger's name includes "Mrs.":

```
titanic_train %>%  
  filter(str_detect(Name, "Mrs\\.\\.\\.")) %>%  
  select(Name) %>%  
  slice(1:3)
```

```
##                               Name  
## 1 Cumings, Mrs. John Bradley (Florence Briggs Thayer)  
## 2      Futrelle, Mrs. Jacques Heath (Lily May Peel)  
## 3 Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)
```

The `str_extract` function can be used to extract a string (if it exists) from each value in a character vector. It follows similar conventions to `str_detect`:

```
## Generic code
```

```
str_extract(string = [vector you want to check],  
            pattern = [pattern you want to check for])
```

Regular expressions

For example, you might want to extract “Mrs.” if it exists in a passenger’s name:

```
titanic_train %>%  
  mutate(mrs = str_extract(Name, "Mrs\\.\\.")) %>%  
  select(Name, mrs) %>%  
  slice(1:3)
```

##	Name	mrs
## 1	Braund, Mr. Owen Harris	<NA>
## 2	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	Mrs.
## 3	Heikkinen, Miss. Laina	<NA>

Notice that now we’re creating a new column (`mrs`) that either has “Mrs.” (if there’s a match) or is missing (`NA`) if there’s not a match.

Regular expressions

For this first example, we were looking for an exact string (“Mrs”). However, you can use patterns that match a particular pattern, but not an exact string. For example, we could expand the regular expression to find “Mr.” or “Mrs.”:

```
titanic_train %>%  
  mutate(title = str_extract(Name, "Mr[s]*\\\\")) %>%  
  select(Name, title) %>%  
  slice(1:3)
```

	Name	title
## 1	Braund, Mr. Owen Harris	Mr.
## 2	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	Mrs.
## 3	Heikkinen, Miss. Laina	<NA>

This pattern uses `[s]*` to match zero or more “s”s at this spot in the pattern.

Regular expressions

In the previous code, we found “Mr.” and “Mrs.”, but missed “Miss.”. We could tweak the pattern again to try to capture that, as well. For all three, we have the pattern that it starts with “M”, has some lowercase letters, and then ends with “.”.

```
titanic_train %>%  
  mutate(title = str_extract(Name, "M[a-z]+\\\\")) %>%  
  select(Name, title) %>%  
  slice(1:3)
```

	Name	title
## 1	Braund, Mr. Owen Harris	Mr.
## 2	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	Mrs.
## 3	Heikkinen, Miss. Laina	Miss.

Regular expressions

Sometimes, you want to match a pattern, but then only subset a part of it. For example, each passenger seems to have a title (“Mr.”, “Mrs.”, etc.) that comes after “,” and before “.”. We can use this pattern to find the title, but then we get some extra stuff with the match:

```
titanic_train %>%  
  mutate(title = str_extract(Name, "[A-Z][a-z]*\\.")) %>%  
  select(title) %>%  
  slice(1:3)
```

```
##      title  
## 1      , Mr.  
## 2    , Mrs.  
## 3    , Miss.
```

Regular expressions

We are getting things like “, Mr. ”, when we really want “Mr”. We can use the `str_match` function to do this. We group what we want to extract from the pattern in parentheses, and then the function returns a matrix. The first column is the full pattern match, and each following column gives just what matches within the groups.

```
head(str_match(titanic_train$Name,  
              pattern = ", ([A-Z] [a-z]*)\\s\\."))
```

```
##      [,1]      [,2]  
## [1,] ", Mr."  "Mr"  
## [2,] ", Mrs." "Mrs"  
## [3,] ", Miss." "Miss"  
## [4,] ", Mrs." "Mrs"  
## [5,] ", Mr."  "Mr"  
## [6,] ", Mr."  "Mr"
```

Regular expressions

To get just the title, then, we can run:

```
titanic_train %>%  
  mutate(title =  
    str_match(Name, " ([A-Z][a-z]*)\\.").[, 2]) %>%  
  select(Name, title) %>%  
  slice(1:3)
```

```
##                               Name title  
## 1                        Braund, Mr. Owen Harris      Mr  
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer)  Mrs  
## 3                        Heikkinen, Miss. Laina      Miss
```

The `[, 2]` pulls out just the second column from the matrix returned by `str_match`.

Regular expressions

Here are some of the most common titles:

```
titanic_train %>%  
  mutate(title =  
    str_match(Name, "([A-Z][a-z]*)\\.", 2)) %>%  
  group_by(title) %>% summarize(n = n()) %>%  
  arrange(desc(n)) %>% slice(1:5)  
  
## `summarise()` ungrouping output (override with `.groups` argu  
  
## # A tibble: 5 x 2  
##   title      n  
##   <chr>  <int>  
## 1 Mr      517  
## 2 Miss    182  
## 3 Mrs     125  
## 4 Master   40  
## 5 Dr       7
```

Regular expressions

The following slides have a few other examples of regular expressions in action with this dataset.

Get just names that start with ("^") the letter "A":

```
titanic_train %>%  
  filter(str_detect(Name, "^A")) %>%  
  select(Name) %>%  
  slice(1:3)
```

##	Name
## 1	Allen, Mr. William Henry
## 2	Andersson, Mr. Anders Johan
## 3	Asplund, Mrs. Carl Oscar (Selma Augusta Emilia Johansson)

Regular expressions

Get names with “II” or “III” ({2,} says to match at least two times):

```
titanic_train %>%  
  filter(str_detect(Name, "I{2,}")) %>%  
  select(Name) %>%  
  slice(1:3)
```

```
##                               Name  
## 1  Carter, Master. William Thornton II  
## 2  Roebling, Mr. Washington Augustus II
```

Regular expressions

Get names with “Andersen” or “Anderson” (alternatives in square brackets):

```
titanic_train %>%  
  filter(str_detect(Name, "Anders[eo]n")) %>%  
  select(Name)
```

```
##                               Name  
## 1 Andersen-Jensen, Miss. Carla Christine Nielsine  
## 2                        Anderson, Mr. Harry  
## 3                        Walker, Mr. William Anderson  
## 4                        Olsvigen, Mr. Thor Anderson  
## 5      Soholt, Mr. Peter Andreas Lauritz Andersen
```

Regular expressions

Get names that start with (“^” outside of brackets) the letters “A” and “B”:

```
titanic_train %>%  
  filter(str_detect(Name, "^[AB]")) %>%  
  select(Name) %>%  
  slice(1:3)
```

```
##              Name  
## 1 Braund, Mr. Owen Harris  
## 2 Allen, Mr. William Henry  
## 3 Bonnell, Miss. Elizabeth
```


Regular expressions

Get names that end with (“\$”) the letter “b” (either lowercase or uppercase):

```
titanic_train %>%  
  filter(str_detect(Name, "[bB]$")) %>%  
  select(Name)
```

```
##                               Name  
## 1   Emir, Mr. Farred Chehab  
## 2 Goldschmidt, Mr. George B  
## 3           Cook, Mr. Jacob  
## 4           Pasic, Mr. Jakob
```

Regular expressions

There is a family of older, base R functions called `grep` that does something very similar.

You may see these functions in example code.

Regular expressions

WHenever I learn a new skill I concoct elaborate fantasy scenarios where it lets me save the day.

OH NO! THE KILLER MUST HAVE FOLLOWED HER ON VACATION!



BUT TO FIND THEM WE'D HAVE TO SEARCH THROUGH 200 MB OF EMAILS LOOKING FOR SOMETHING FORMATTED LIKE AN ADDRESS!



IT'S HOPELESS!

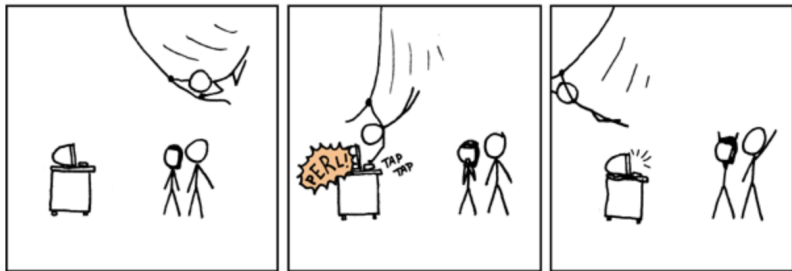
Source: xkcd

Regular expressions



Source: xkcd

Regular expressions



Source: xkcd

Regular expressions

For more on these patterns, see:

- Help file for the `stringi-search-regex` function in the `stringi` package (which should install when you install `stringr`)
- Chapter 14 of R For Data Science
- <http://gskinner.com/RegExr>: Interactive tool for helping you build regular expression pattern strings

Tidy select

There are `tidyverse` functions to make selecting variables more straightforward. You can call these functions as arguments of the `select` function to streamline variable selection. Examples include: `starts_with()`, `ends_with()`, and `contains()`.

Tidy select (helpers)

Here we use `starts_with("t")` to select all variables that begin with t.

```
titanic_train %>%  
  select(starts_with("t")) %>%  
  slice(1:3)
```

```
##           Ticket  
## 1      A/5 21171  
## 2      PC 17599  
## 3 STON/O2. 3101282
```


Tidy select

There are also tidyverse functions that allow us to easily operate on a selection of variables. These functions are called `scoped variants`. You can identify these functions by these `_all`, `_at`, and `_if` suffixes.

Tidy select (*_if)

Here we use `select_if` to select all the numeric variables in a dataframe and covert their names to lower case (a handy function to tidy the variable names).

```
titanic_train %>%  
  select_if(is.numeric, tolower) %>%  
  slice(1:3)
```

##	passengerid	survived	pclass	age	sibsp	parch	fare
## 1	1	0	3	22	1	0	7.2500
## 2	2	1	1	38	1	0	71.2833
## 3	3	1	3	26	0	0	7.9250

Tidy select (*_if)

The `select_if` function takes the following form.

Generic code

```
new_df <- select_if(old_df,  
                    .predicate [selects the variable to keep],  
                    .funcs = [the function to apply to  
                              the selected column names])
```

Tidy select (*_at)

Here we use `select_at` to select all the variables that contain `ss` in their name and then convert their names to lower case (a handy function to tidy the variable names).

```
titanic_train %>%  
  select_at(vars(contains("ss")), tolower) %>%  
  slice(1:3)
```

```
##   passengerid pclass  
## 1           1      3  
## 2           2      1  
## 3           3      3
```